

Run SQL data queries with pandas

How-to: Run SQL data queries with pandas

Use pandas to do joins, grouping, aggregations, and analytics on datasets in Python.

By Yuli Vasiliev | *March 2021*



Python's pandas library, with its fast and flexible data structures, has become the de facto standard for data-centric Python applications, offering a rich set of built-in facilities to analyze details of structured data. Built on top of other core Python libraries, such as NumPy, SQLAlchemy, and Matplotlib, pandas leverages these libraries behind the scenes for quick and easy data manipulations, allowing you to take advantage of their functionality with less coding. For example, the `read_sql()` and `to_sql()` pandas methods use SQLAlchemy under the hood, providing a unified way to send pandas data in and out of a SQL database.

This article illustrates how you can use pandas to combine datasets, as well as how to group, aggregate, and analyze data in them. For comparison purposes, you'll also see how these same tasks can be addressed with SQL.

Creating database structures for article examples

To follow along with the examples in this article, you need to create several example tables in an Oracle database by executing the `pandas_article.sql` script that accompanies the article. Also make sure you have the pandas, SQLAlchemy, and cx_Oracle libraries installed in your Python environment. You can install them using the `pip` command:

```
pip install pandas
pip install SQLAlchemy
pip install cx_Oracle
```

[Copy](#)

For details on how to install pandas, refer to the [documentation](#). For SQLAlchemy installation details, refer to the [SQLAlchemy documentation](#). For details on how to install cx_Oracle, refer to the [cx_Oracle Installation page](#). You might also want to look at the [cx_Oracle Initialization page](#).

Loading data from Oracle Database to pandas DataFrames

After executing the `pandas_article.sql` script, you should have the orders and details database tables populated with example data. The following script connects to the database and loads the data from the orders and details tables into two separate DataFrames (in pandas, DataFrame is a key data structure designed to work with tabular data):

```
import pandas as pd
import cx_Oracle
import sqlalchemy
from sqlalchemy.exc import SQLAlchemyError
try:
    engine = sqlalchemy.create_engine("oracle+cx_oracle://usr:pswd@localhost/?serv
orders_sql = """SELECT * FROM orders""";
df_orders = pd.read_sql(orders_sql, engine)
details_sql = """SELECT * FROM details""";
df_details = pd.read_sql(details_sql, engine)
print(df_orders)
print(df_details)
except SQLAlchemyError as e:
    print(e)
```

In this example, you use `sqlalchemy` to create an engine to connect to an Oracle database. Using a SQLAlchemy engine allows you to pass in the `arraysize` argument that will be used when `cx_Oracle.Cursor` objects are created.

The `arraysize` attribute of the `cx_Oracle.Cursor` object is used to tune the number of rows internally fetched and buffered when fetching rows from `SELECT` statements and `REF CURSOR`. By default, this attribute is set to 100, which is perfectly acceptable when you need to load a small amount of data from the database. However, when you're dealing with large amounts of data, you should increase the value of `arraysize` to reduce the number of round trips between your script and the database and, therefore, improve performance. For further details on how you can use `arraysize`, refer to the [Tuning cx_Oracle documentation page](#).

The script prints the `df_orders` and `df_details` DataFrames loaded from the database, producing the following output:

	PONO	ORDATE	EMPL
0	7723510	2020-12-15	John Holland
1	5626234	2020-12-15	Tim Lewis
2	7723533	2020-12-15	John Holland
3	7823675	2020-12-16	Maya Candy
4	5626376	2020-12-16	Tim Lewis
5	5626414	2020-12-17	Dan West
6	7823787	2020-12-17	Maya Candy
7	5626491	2020-12-17	Dan West

	PONO	LINEID	ITEM	BRAND	PRICE	QUANTITY	DISCOUNT
0	7723510	1	Swim Shorts	Hurley	17.95	1	0
1	7723510	2	Jacket	Oakley	142.33	1	0
2	5626234	1	Socks	Vans	16.15	4	15
3	7723533	1	Jeans	Quiksilver	84.90	2	25
4	7723533	2	Socks	Mons Royale	10.90	2	0
5	7723533	3	Socks	Stance	12.85	2	20
6	7823675	1	T-shirt	Patagonia	35.50	3	0
7	5626376	1	Hoody	Animal	44.05	1	0
8	5626376	2	Cargo Shorts	Animal	38.60	1	12
9	5626414	1	Shirt	Volcom	78.55	2	0
10	7823787	1	Boxer Shorts	Superdry	30.45	2	18
11	7823787	2	Shorts	Barts	35.90	1	0
12	5626491	1	Cargo Shorts	Billabong	48.74	1	22
13	5626491	2	Sweater	Dickies	65.95	1	0

 Copy

Joining DataFrames

The `DataFrame.merge()` method is designed to address this task for two DataFrames. The method allows you to explicitly specify columns in the DataFrames, on which you want to join those DataFrames. You can also specify the type of join to produce the desired result set. By default, `merge()` creates an inner join on the column that the DataFrames being joined have in common. So, you can join the `df_orders` and `df_details` DataFrames created in the previous section with the following simple call of `merge()`:

```
df_orders_details = df_orders.merge(df_details)
```

 Copy

If you print the `df_orders_details` DataFrame, it should look as follows:

	PONO	ORDATE	EMPL	LINEID	ITEM	BRAND	PRI
0	7723510	2020-12-15	John Holland	1	Swim Shorts	Hurley	17.
1	7723510	2020-12-15	John Holland	2	Jacket	Oakley	142
2	5626234	2020-12-15	Tim Lewis	1	Socks	Vans	16.
3	7723533	2020-12-15	John Holland	1	Jeans	Quiksilver	84.
4	7723533	2020-12-15	John Holland	2	Socks	Mons Royale	10.
5	7723533	2020-12-15	John Holland	3	Socks	Stance	12.
6	7823675	2020-12-16	Maya Candy	1	T-shirt	Patagonia	35.
7	5626376	2020-12-16	Tim Lewis	1	Hoody	Animal	44.
8	5626376	2020-12-16	Tim Lewis	2	Cargo Shorts	Animal	38.
9	5626414	2020-12-17	Dan West	1	Shirt	Volcom	78.
10	7823787	2020-12-17	Maya Candy	1	Boxer Shorts	Superdry	30.
11	7823787	2020-12-17	Maya Candy	2	Shorts	Barts	35.
12	5626491	2020-12-17	Dan West	1	Cargo Shorts	Billabong	48.
13	5626491	2020-12-17	Dan West	2	Sweater	Dickies	65.

 Copy

After joining two datasets into a single one, you may still need to modify it before you can perform analysis. In the case of `df_orders_details` being discussed here, you might need to add some new columns, calculating their values based on the values in the existing columns. Thus, you might need to add a `TOTAL` column that contains the extended item price (price multiplied by quantity and minus discount), for example:

```
df_orders_details['TOTAL'] = df_orders_details.PRICE * df_orders_details.QUANTITY
```

 Copy

```
df_orders_details[['OFF']] = df_orders_details[['TOTAL', 'OFF']].round(2)
```

 Copy

Since all the float columns in the `df_orders_details` DataFrame contain monetary values, you can specify two decimal places to round each float column to

```
df_orders_details = df_orders_details.round(2)
```

 Copy

Some columns in the DataFrame may not be needed for the analysis you want to perform. So, you can keep only those columns that are needed. In the `df_orders_details` DataFrame, for example, if you want to group sales data (both totals and discounts) by order dates and employees, you can keep just these four columns:

```
df_sales = df_orders_details[['ORDATE', 'EMPL', 'TOTAL', 'OFF']]
```

 Copy

If you print the DataFrame, it will look like this:

	ORDATE	EMPL	TOTAL	OFF
0	2020-12-15	John Holland	17.95	0.00
1	2020-12-15	John Holland	142.33	0.00
2	2020-12-15	Tim Lewis	54.91	9.69
3	2020-12-15	John Holland	127.35	42.45
4	2020-12-15	John Holland	21.80	0.00
5	2020-12-15	John Holland	20.56	5.14
6	2020-12-16	Maya Candy	106.50	0.00
7	2020-12-16	Tim Lewis	44.05	0.00
8	2020-12-16	Tim Lewis	33.97	4.63
9	2020-12-17	Dan West	157.10	0.00
10	2020-12-17	Maya Candy	49.94	10.96
11	2020-12-17	Maya Candy	35.90	0.00
12	2020-12-17	Dan West	38.02	10.72
13	2020-12-17	Dan West	65.95	0.00

 Copy

```
CREATE VIEW sales_v AS
SELECT
    ordate,
    empl,
    price*quantity*(1-discount/100) AS total,
    price*quantity*(discount/100) AS off
FROM orders INNER JOIN details
ON orders.pono = details.pono;
```

 Copy

Grouping and aggregating data

Using the `DataFrame.groupby()` method you can split a `DataFrame`'s data into subsets (groups) that have matching values for one or more columns, and then apply an aggregate function to each group. In the following example, you group by the `ORDATE` and `EMPL` columns in the `df_sales` `DataFrame` and then apply the `sum()` aggregate function to the `TOTAL` and `OFF` columns within the formed groups:

```
df_date_empl = df_sales.groupby(['ORDATE', 'EMPL']).sum()
```

 Copy

The generated `DataFrame` should look as shown below:

ORDATE	EMPL	TOTAL	OFF
2020-12-15	John Holland	329.99	47.59
	Tim Lewis	54.91	9.69
2020-12-16	Maya Candy	106.50	0.00
	Tim Lewis	78.02	4.63
2020-12-17	Dan West	261.07	10.72
	Maya Candy	85.84	10.96

 Copy

One problem here is that the aggregate function you apply to the `groupby` object is applied to each numeric column of the `DataFrame`. But what if you need to apply multiple aggregate functions to multiple `groupby`

```
df_aggs = df_sales.groupby(['ORDATE', 'EMPL']).agg({'TOTAL': ['sum', 'mean'], 'OFF
```

 Copy

The above example illustrates how you can select a certain column for aggregation and perform different aggregations per column. If you print `df_aggs`, it will look as follows:

ORDATE	EMPL	TOTAL		OFF
		sum	mean	max
2020-12-15	John Holland	329.99	66.00	42.45
	Tim Lewis	54.91	54.91	9.69
2020-12-16	Maya Candy	106.50	106.50	0.00
	Tim Lewis	78.02	39.01	4.63
2020-12-17	Dan West	261.07	87.02	10.72
	Maya Candy	85.84	42.92	10.96

 Copy

You might want to flatten a hierarchical index in columns. This can be done as follows:

```
df_aggs.columns = df_aggs.columns.map('_'.join).str.strip()
```

 Copy

This will change the column names as shown below:

TOTAL_sum	TOTAL_mean	OFF_max
ORDATE	EMPL	
...		

 Copy

To generate the same result set with a query to the article sample database, you could issue the following `SELECT` statement against the `sales_v` view that you should have created previously:

```
    OFF_max,  
    ROUND(SUM(total),2) TOTAL_sum,  
    ROUND(AVG(total),2) TOTAL_mean,  
    ROUND(MAX(off),2) OFF_max  
FROM  
    sales_v  
GROUP BY  
    ordate, empl  
ORDER BY  
    ordate;
```

 Copy

Analytical processing within groups of data

In practice, you may not always need to view data in summarized format, aggregating a group of rows into a single resulting row as illustrated in the previous example. In contrast, you may need to do some analytical processing within a group of rows so the number of rows in the group remains the same. For example, if you want to compare the salary of each employee in a department with the average salary of the employees in this department, this processing does not imply any reduction in the number of rows in the dataset—the number of rows must match the number of employees, both before and after processing.

Let's illustrate this analytical processing with a second, more complex example. Imagine you want to analyze stock price data for a list of tickers over a certain period of time. To start, you want to weed out the tickers whose prices dropped below 1% of the previous day's price over the period. To accomplish this, you need to group data by ticker symbol, ordering the rows by date in each group. Then you can iterate over the rows in a group, comparing the stock price in the current row with the price in the previous row. If the price in a current row is less than the price in the previous row by more than 1%, then the entire group of rows must be excluded from the result set. This section describes how you could implement this filtering.

The example uses stock data obtained via the `yfinance` library, a Python wrapper for the Yahoo Finance API, which you can install with the `pip` command, as follows:

```
pip install yfinance
```

 Copy

In the following script, you get stock data for several popular stocks for a five-day period:

```
import pandas as pd  
import yfinance as yf
```



```
hist = tkr.history(period='5d')
hist['Symbol']=ticker
stocks = stocks.append(hist[['Symbol', 'Close']].rename(columns={'Close': 'Price'}
```

 Copy

yfinance returns a requested dataset as a pandas DataFrame with the `Date` column as the index. Assuming you are targeting the closing prices only, you keep only the `Symbol` and `Close` columns, having renamed the latter to `Price` for comprehension. As a result, the data in the DataFrame might look like this:

Date	Symbol	Price
2020-12-18	AAPL	126.660004
2020-12-21	AAPL	128.229996
2020-12-22	AAPL	131.880005
2020-12-23	AAPL	130.960007
2020-12-24	AAPL	131.970001
2020-12-18	TSLA	695.000000
2020-12-21	TSLA	649.859985
2020-12-22	TSLA	640.340027
2020-12-23	TSLA	645.979980
2020-12-24	TSLA	661.770020
2020-12-18	FB	276.399994
2020-12-21	FB	272.790009
2020-12-22	FB	267.089996
2020-12-23	FB	268.109985
2020-12-24	FB	267.399994
2020-12-18	ORCL	65.059998
2020-12-21	ORCL	64.480003
2020-12-22	ORCL	65.150002
2020-12-23	ORCL	65.300003
2020-12-24	ORCL	64.959999
2020-12-18	AMZN	3201.649902
2020-12-21	AMZN	3206.179932
2020-12-22	AMZN	3206.520020
2020-12-23	AMZN	3185.270020
2020-12-24	AMZN	3172.689941

 Copy

From the above row set, you need to select the rows related to only those symbols whose prices did not drop below 1% of the previous day's price. For this, you need a mechanism that will allow you to compare the `Price` value of a row with the `Price` value of the previous row within a symbol group. The following line of code

```
stocks['Prev'] = stocks.groupby(['Symbol'])['Price'].shift(1)
```

 Copy

The updated row set should look as follows:

Date	Symbol	Price	Prev
2020-12-18	AAPL	126.660004	NaN
2020-12-21	AAPL	128.229996	126.660004
2020-12-22	AAPL	131.880005	128.229996
2020-12-23	AAPL	130.960007	131.880005
2020-12-24	AAPL	131.970001	130.960007
2020-12-18	TSLA	695.000000	NaN
2020-12-21	TSLA	649.859985	695.000000
2020-12-22	TSLA	640.340027	649.859985
2020-12-23	TSLA	645.979980	640.340027
2020-12-24	TSLA	661.770020	645.979980
2020-12-18	FB	276.399994	NaN
2020-12-21	FB	272.790009	276.399994
2020-12-22	FB	267.089996	272.790009
2020-12-23	FB	268.109985	267.089996
2020-12-24	FB	267.399994	268.109985
2020-12-18	ORCL	65.059998	NaN
2020-12-21	ORCL	64.480003	65.059998
2020-12-22	ORCL	65.150002	64.480003
2020-12-23	ORCL	65.300003	65.150002
2020-12-24	ORCL	64.959999	65.300003
2020-12-18	AMZN	3201.649902	NaN
2020-12-21	AMZN	3206.179932	3201.649902
2020-12-22	AMZN	3206.520020	3206.179932
2020-12-23	AMZN	3185.270020	3206.520020
2020-12-24	AMZN	3172.689941	3185.270020

 Copy

The `Prev` results for the first day of the observation period are `NaN` because this example does not track what happened before the five-day range.

Now you can find those rows where the ratio of the price to the previous price is less than 99%, for example:

```
stocks_to_exclude = stocks[stocks['Price']/stocks['Prev'] < .99]
```

Date	Symbol	Price	Volume
2020-12-21	TSLA	649.859985	695.000000
2020-12-22	TSLA	640.340027	649.859985
2020-12-21	FB	272.790009	276.399994
2020-12-22	FB	267.089996	272.790009

 Copy

You can extract the symbols presented in the above rows as follows:

```
exclude_list = list(set(stocks_to_exclude['Symbol'].tolist()))
```

 Copy

Here you extract the values of the `Symbol` column in the `stocks_to_exclude` DataFrame, converting those values to a list. To exclude duplicates from this list, you convert it into a set and then back to a list (one of the most popular ways to remove duplicates from a list).

```
['TSLA', 'FB']
```

 Copy

Next you need to exclude those rows from the `stocks` DataFrame that includes the above names in the `Symbol` field. This can be implemented as the following one-liner:

```
stocks_filtered = stocks[~stocks['Symbol'].isin(exclude_list)][['Symbol', 'Price']
```

 Copy

You pass in the `exclude_list` that you created previously to the `stocks['Symbol'].isin()` function to get a Series of booleans indicating if each value in the `stocks['Symbol']` Series is in `exclude_list`. You put a tilde sign in front of `stocks['Symbol']` to invert the boolean Series returned by `isin()`. You pass in this inverted boolean Series to the `[]` operator of the `stocks` DataFrame to return all the rows that do not contain in the `Symbol` column symbols found in the `exclude_list`. So, the resulting DataFrame should look as follows:

```

2020-12-18  AAPL  128.000001
2020-12-21  AAPL  128.229996
2020-12-22  AAPL  131.880005
2020-12-23  AAPL  130.960007
2020-12-24  AAPL  131.970001
2020-12-18  ORCL  65.059998
2020-12-21  ORCL  64.480003
2020-12-22  ORCL  65.150002
2020-12-23  ORCL  65.300003
2020-12-24  ORCL  64.959999
2020-12-18  AMZN  3201.649902
2020-12-21  AMZN  3206.179932
2020-12-22  AMZN  3206.520020
2020-12-23  AMZN  3185.270020
2020-12-24  AMZN  3172.689941

```

 Copy

With the help of analytical SQL, you can get the same result set with a single query to the database. Before you can do this, however, you need to save the unfiltered row set to the article database, which should contain the `stocks` table for storing this data (refer back to the `pandas_article.sql` script that you should have run at the very beginning).

To conform to the structure of the `stocks` database table, you need to modify the `stocks` DataFrame that contains the unfiltered data of this example. This can be done with the following lines of code:

```

stocks_to_db = stocks[['Symbol', 'Price']].reset_index().rename(columns={'Date':
stocks_to_db = stocks_to_db.astype({'Dt': str})

```

 Copy

In the first line, you specify the columns to include in the result set: `Symbol` and `Price`. By resetting the index, you add `Date` to this column list. To conform to the name of this column in the `stocks` database table, you rename it to `Dt`. The `round(2)` function rounds the values in the `Price` column to two decimal places. In the second line, you cast the `Dt` column to the `str` type, because pandas sets it to datetime by default.

Finally, you need to convert the `stocks_to_db` DataFrame to a structure that is passable to a method that can do a bulk insert operation. In the following line of code, you convert the `stocks_to_db` DataFrame to a list of tuples:

```

data = list(stocks_to_db.itertuples(index=False, name=None))

```

The following script uses the above list of tuples to upload the data it contains to the database. Storing the data you work with can be useful when you're going to reuse it.

```
import cx_Oracle
try:
    conn = cx_Oracle.connect("usr", "pswd", "localhost/orcl")
    cursor = conn.cursor()
    #defining the query
    query_add_stocks = """INSERT INTO stocks (dt, symbol, price)
                           VALUES (TO_DATE(:1, 'YYYY-MM-DD'), :2, :3)"""
    #inserting the stock rows
    cursor.executemany(query_add_stocks, data)
    conn.commit()
except cx_Oracle.DatabaseError as exc:
    err = exc.args
    print("Oracle-Error-Code:", err.code)
    print("Oracle-Error-Message:", err.message)
finally:
    cursor.close()
    conn.close()
```

 Copy

In this script, you connect to the database and obtain a cursor object to interact with it. You use the `cursor.executemany()` method that inserts all the rows from the data list of tuples into the database in a single round trip.

After the successful execution of the above script, you can issue queries against the stocks table. To get the row set you had in the `stocks_filtered` DataFrame, you can issue the following query:

```
SELECT s.* FROM stocks s
LEFT JOIN
(SELECT DISTINCT(symbol) FROM
 (SELECT price/LAG(price) OVER (PARTITION BY symbol ORDER BY dt) AS dif, symbol
  ON a.symbol = s.symbol WHERE a.symbol IS NULL;
```

 Copy

Conclusion

reshaping original datasets as needed.

Dig deeper

- [Learn more about how to install pandas.](#)
- [Learn more about how to install SQLAlchemy.](#)
- [Learn more about how to install cx_Oracle.](#)
- [Read the cx_Oracle Initialization guide.](#)
- [Get the sample dataset for this article.](#)

Illustration: Wes Rowell

Yuli Vasiliev is a programmer, freelance writer, and consultant. He is the author of [Natural Language Processing with Python and spaCy](#).

Connect



Contact us



Subscribe



Share this page

Resources for

Careers
Developers
Investors
Partners
Startups
Students and Educators

Why Oracle

Analyst Reports
Gartner MQ for ERP Cloud
Cloud Economics
Corporate Responsibility
Diversity and Inclusion
Security Practices

Learn

What is cloud computing?
What is CRM?
What is Docker?
What is Kubernetes?
What is Python?
What is SaaS?

What's New

Try Oracle Cloud Free Tier
Oracle Arm Processors
Oracle and Premier League
Oracle and Red Bull Racing Honda
Employee Experience Platform
Oracle Support Rewards

Contact Us

US Sales: +1.800.633.0738
How can we help?
Subscribe to emails
Events
News
Blogs

